

Meta Title

RAG Development Services for Enterprise AI | WildMind AI

Meta Description

Custom RAG development services that connect LLMs to your business data. Build secure, scalable AI with accurate retrieval, lower hallucinations, and enterprise deployment.

URL: <https://solutions.wildmindai.com/services/rag-development>

Full Page Content

Eyebrow: ENTERPRISE AI ENGINEERING

RAG development services grounded in your own data, not a model's guesswork.

WildMind AI Solutions delivers custom RAG development services that connect large language models (LLMs) to your business data, enabling AI to generate accurate, context-aware, and trustworthy responses. We build secure, scalable Retrieval-Augmented Generation (RAG) systems that transform enterprise knowledge into intelligent search, automation, and decision-making tools, helping organizations improve productivity while reducing AI hallucinations.

[Book a Strategy Call] [View Our Work]

Trusted by engineering teams building production AI

Eyebrow text: Core Capabilities

Our RAG Development Capabilities

Custom RAG Development

We design retrieval-augmented generation systems around how your data is actually structured, not a generic template. That includes chunking strategy, retrieval method, reranking, and how retrieved context gets assembled into the final prompt, tuned to your document types and query patterns.

AI Agent Development and Workflow Automation

Beyond single-pass retrieval, our AI agent development practice builds agentic RAG systems that decompose complex queries, run multi-step retrieval across several data sources, call

internal tools and APIs, and maintain context across a multi-turn task. We build these on frameworks like LangChain and LlamaIndex, orchestrating retrieval, tool calls, and reasoning steps into AI workflow automation that goes beyond "answer a question" into "complete a task."

Vector Database Implementation

We implement and tune vector databases including Pinecone, Weaviate, Milvus, Qdrant, ChromaDB, and FAISS, choosing the right one based on your scale, latency requirements, metadata filtering needs, and hosting constraints (cloud, VPC, or on-prem).

Semantic, Keyword, and Hybrid Search

We build hybrid retrieval that combines dense semantic search with traditional keyword matching and metadata filters, plus reranking layers that push the most relevant passages to the top before generation, the detail that separates a demo from something a compliance team will sign off on.

LLM Integration Across Providers

We integrate GPT, Claude, Gemini, Llama, and open-source models with your retrieval pipeline, and we're not locked into one vendor, model choice is driven by your latency, cost, data residency, and accuracy requirements, with the retrieval layer built to be swapped between models with minimal rework.

AI Chatbot Development and Conversational AI

Deliver accurate, context-aware conversations with RAG-powered AI applications. We build AI chatbots, conversational AI assistants, and document search solutions that retrieve information from your latest business data, providing reliable answers backed by verifiable source references.

Secure, Compliant Enterprise Deployment

Protect sensitive business data with secure RAG deployments designed for enterprise environments. We implement role-based access control, encrypted data handling, audit logging, and deploy within your own VPC or on-premises infrastructure to meet security and data residency requirements.

Mid Page CTA Section Block

Not Sure If Retrieval Is Your Actual Bottleneck?

Many AI performance issues start with retrieval, not the model. WildMind AI Solutions assesses your RAG pipeline to identify bottlenecks and improve response accuracy.

Get a Free RAG Assessment

Our RAG Development Process

01 Discovery and Use Case Assessment

We map your business goals, data sources, users, and compliance constraints to identify the highest-value RAG use case to build first, not the most technically interesting one.

02 Data Preparation and Knowledge Engineering

We clean, structure, and chunk source content from PDFs, wikis, cloud storage, CRMs, ERPs, and internal databases, with a chunking strategy matched to how that content is actually written and queried.

03 Embedding and Vector Database Setup

We generate embeddings and stand up a vector database sized and configured for your query volume, latency targets, and metadata filtering needs.

04 Retrieval Pipeline Development

We build the retrieval logic, hybrid search, reranking, metadata filters, query decomposition for complex or multi-part questions, that determines whether the right context reaches the model in the first place.

05 LLM Integration and Prompt Engineering

We connect the retrieval layer to your chosen LLM(s), engineering the prompt assembly and response format so answers stay grounded in retrieved context and cite their sources.

06 Evaluation, Deployment, and Monitoring

Before launch, we test retrieval accuracy and hallucination rate against real queries, then deploy with monitoring in place so retrieval quality is tracked continuously, not assumed.

Eyebrow text: Why WildMind AI Solutions

Why Choose WildMind AI Solutions for RAG Development Services

Retrieval-First Diagnosis

Before recommending a rebuild, we audit the existing pipeline against the four common failure points. Most stalled pilots need two fixes, not a full restart.

Takes Over Stalled Builds

We regularly pick up RAG projects that stalled with another vendor or an internal team and get them to production without starting from zero.

Named Senior Engineers

You know who is architecting and building your retrieval pipeline before the engagement starts, not after the kickoff call.

Full Source Code Ownership

Every chunking script, prompt template, and eval suite produced during the build is yours at delivery, no exceptions.

Discovery Before Any Build

Requirements, data audit, and architecture get documented before a single line of retrieval code gets written.

Post-Launch SLA Support

L1, L2, and L3 support options are available after launch, so a model upgrade or index migration doesn't become your problem alone.

Case studies

Blogs

FAQ Section Block

1. What are RAG Development Services?

RAG Development Services combine Large Language Models (LLMs) with your business data to deliver accurate, context-aware AI responses. Instead of relying only on pre-trained knowledge, RAG retrieves relevant information from your documents and databases in real time.

2. How is RAG different from fine-tuning an LLM?

RAG improves AI by retrieving up-to-date information from your knowledge base during each query, while fine-tuning changes how a model behaves or responds. Many enterprise AI solutions use both approaches together for the best results.

3. Can you integrate a RAG solution with our existing systems?

Yes. We integrate RAG applications with CRMs, ERPs, cloud storage, websites, internal portals, customer support platforms, and other enterprise systems to ensure seamless access to your business knowledge.

4. Can our RAG solution be deployed in a private cloud or on-premises?

Absolutely. We build secure RAG solutions that can be deployed in your own cloud, private VPC, or on-premises infrastructure with role-based access controls, encryption, and compliance-focused security.

5. How long does it take to build a RAG solution?

The timeline depends on your data sources, integrations, and project complexity. Most RAG MVPs can be delivered within 6–10 weeks, while enterprise-scale implementations typically require additional time for testing, integrations, and optimization.

Black CTA Banner Section

Ready to Fix or Build Your RAG System?

Whether you're starting from scratch or improving an existing solution, WildMind AI Solutions helps you design, develop, and deploy RAG systems that deliver accurate, scalable, and enterprise-ready AI experiences.

[Start Your RAG Transformation Today]